

# Model Context Protocol (MCP): A Comprehensive Technical Review

---

**Report Date:** 2025-07-15

## Executive Summary

The Model Context Protocol (MCP) is an open standard, initiated by Anthropic, designed to create a universal interface between Large Language Models (LLMs) and external data sources, tools, and services. Conceived as a “USB-C port for AI,” MCP addresses the critical challenge of fragmented, ad-hoc integrations that have historically hindered the scalability, security, and contextual awareness of AI applications. This report provides a comprehensive literature review of MCP for a technical audience, synthesizing information from official specifications, academic research, security analyses, and implementation guides. It covers the protocol’s core architecture, its key features, the significant security landscape it navigates, practical implementation guidelines, and a comparative analysis against traditional and contemporary protocols. The findings indicate that MCP represents a pivotal advancement in creating more capable, secure, and interoperable agentic AI systems. However, its power and flexibility introduce a complex threat surface that demands a rigorous, security-first approach from all implementers. The protocol is evolving rapidly, with a strong community focus on standardizing security practices, particularly around authentication and authorization, to ensure its sustainable and safe adoption across the industry.

---

## 1. Introduction to the Model Context Protocol

---

The rapid proliferation of Large Language Models has exposed a fundamental bottleneck in their real-world application: the difficulty of providing them with timely, relevant, and secure access to external context. Before the advent of a standardized protocol, connecting an AI model to a new data source—be it a local file system, an enterprise database, or a third-party API—required bespoke, custom-coded integrations. This approach resulted in a brittle and inefficient ecosystem, characterized by duplicated development effort, inconsistent security controls, and a complex “N times M” integration problem where every AI client needed a unique connector for every tool or server.

To address these challenges, Anthropic introduced and open-sourced the Model Context Protocol (MCP) in late 2024. MCP provides a single, open, and universal standard for AI systems to communicate with the vast landscape of digital information and tools. The protocol’s vision is to foster fungibility between AI clients and servers, breaking down data silos and enabling the creation of complex, composable workflows built upon LLMs. By standardizing how contextual information is shared, how tools are exposed, and how AI-driven actions are executed, MCP aims to unlock the next generation of context-aware, agentic AI applications.

The protocol is designed as a collaborative, open-source project, with Anthropic leading its development and encouraging contributions from the wider community, including major industry players like Google, Microsoft, and Shopify. This collaborative approach is crucial for establishing MCP as a durable industry standard, analogous to the role HTTP played for the web or the Language Server Protocol (LSP) played for developer tools. Its adoption promises to simplify development, enhance security, and ultimately enable AI models to produce more accurate, relevant, and useful responses by grounding them in the specific context of a given task.

## 2. Core Architecture and Protocol Specification

---

The Model Context Protocol is built upon a layered host-client-server architecture and utilizes JSON-RPC 2.0 for structured, stateful communication. This design provides a clear separation of concerns, enhances security through isolation, and supports the progressive addition of features. The official specification, with a current version of 2025-06-18, defines the authoritative requirements for all implementations.

The architecture consists of three primary components. The **MCP Host** is the end-user application, such as an IDE or a desktop AI assistant like Claude Desktop, that serves as the container and coordinator. The host is responsible for managing the lifecycle of client connections, enforcing security policies, handling user consent and authorization, and aggregating context from various sources for the LLM. Within the host, **MCP Clients** are the protocol-aware components that establish and maintain isolated, stateful, one-to-one connections with servers. Each client handles protocol negotiation, message routing, and subscriptions for a single server, ensuring that servers cannot “see into” each other. Finally, **MCP Servers** are lightweight programs or services that act as wrappers or intermediaries, exposing specific capabilities, data, and tools from external systems like databases, APIs, or local file systems through the standardized protocol.

MCP defines several fundamental primitives, or features, that servers can expose and clients can consume. These are negotiated during the connection initialization. **Tools** are executable functions that allow the model to perform actions or retrieve information. They are considered “model-controlled,” meaning the LLM can discover and decide to invoke them based on its understanding of the user’s intent. **Resources** represent structured data or content, such as files or database schemas, identified by a unique URI. They are “application-driven,” meaning the host application determines how this contextual information is used, whether through UI elements, search, or automatic inclusion in a prompt. **Prompts** are pre-defined templates or instructions that guide user interaction, often surfaced as slash commands in a client’s UI. They are “user-controlled,” requiring explicit user initiation.

In addition to these server-exposed features, the protocol defines client-side capabilities. **Sampling** is a powerful feature that allows a server to request an LLM completion from the client. This enables servers to implement complex, agentic behaviors by nesting LLM calls within their own logic, while the client retains full control over model access, permissions, and user approval. **Roots** allow a server to understand its operational boundaries within a file system by querying the client for a list of accessible `file://` URI roots. This architectural foundation, with its clear roles and extensible feature set, is designed to be highly composable, allowing for the creation of layered and hierarchical agent systems where a single component can act as both a server to a primary agent and a client to other specialized services.

## 3. Security Analysis and Threat Landscape

---

The power and flexibility of the Model Context Protocol inherently introduce a significant and complex security landscape. Because MCP grants AI models programmatic access to data and the ability to execute arbitrary code, its implementation demands a rigorous, defense-in-depth security posture. The protocol’s design principles—such as server isolation and mandatory user consent—provide a foundation for security, but the responsibility for robust implementation falls on the developers of hosts, clients, and servers.

A primary category of threats involves the manipulation of the LLM through malicious context provided by a compromised or rogue MCP server. **Tool Poisoning**, also known as Tool Description Prompt Injection,

tion, is a critical vulnerability where an attacker embeds malicious instructions within the descriptive metadata of a tool. These instructions, invisible to the user, can persuade the LLM to execute unintended actions, such as exfiltrating sensitive local files. Researchers have demonstrated exploits where a simple calculator tool’s description contained hidden directives to find and upload SSH keys. A related vulnerability is the **“Rug Pull” attack**, where a server initially presents benign tools but dynamically redefines their functionality to become malicious after the user has granted trust.

Server-side vulnerabilities in the code of MCP servers themselves present another major risk. Due to the nature of their function, many servers execute system commands or interact with file systems based on parameters from the LLM. Without stringent input sanitization, this can lead to classic vulnerabilities like **Command Injection** (resulting in Remote Code Execution), **Path Traversal** (allowing unauthorized file access), and **Server-Side Request Forgery (SSRF)**. Research has shown that a significant percentage of publicly available MCP server implementations contain these types of unsafe coding practices.

The protocol’s architecture also creates avenues for more sophisticated attacks. The **Confused Deputy Problem** can arise when an MCP server acts as a proxy to a third-party API. An attacker can trick the proxy server into using its legitimate credentials to perform unauthorized actions on the user’s behalf. **Session Hijacking** is another risk where an attacker who obtains a session ID could potentially impersonate a legitimate client. Furthermore, the composability of MCP can be exploited through **Composability Chaining**, where a seemingly benign server proxies requests to a hidden, malicious server, which injects tainted data or instructions back into the communication flow.

Mitigation strategies are multi-layered and are a central focus of the protocol’s ongoing evolution. A pivotal development is the shift away from servers acting as their own OAuth providers towards mandatory integration with external Identity Providers (IdPs). This delegates complex authentication and authorization logic to battle-tested platforms. The specification now promotes strong, scope-based access control, mandating the use of modern standards like OAuth 2.1 with PKCE (Proof Key for Code Exchange). For implementers, best practices are paramount: treat all AI-provided input as untrusted, apply rigorous input validation and sanitization, run servers in sandboxed environments with least-privilege principles, and implement comprehensive monitoring and logging to detect anomalous activity. For users, the core defense remains vigilance: only connect to trusted, verified MCP servers.

## 4. Implementation and Best Practices

---

Implementing a robust and secure Model Context Protocol integration requires adherence to established best practices covering environment setup, server development, performance optimization, and context management. The ecosystem provides official Software Development Kits (SDKs) for numerous languages, including Python, TypeScript, Java, and C#, which significantly simplify the development process.

The initial setup involves preparing a suitable development environment. This includes ensuring adequate hardware resources (e.g., multi-core CPU, sufficient RAM for context caching) and installing the necessary software dependencies. A dedicated virtual environment is crucial for isolating project dependencies and ensuring stability; modern package managers like `uv` are recommended for their speed and efficiency. API keys and other secrets must be managed securely using environment variables stored in a `.env` file, which should be excluded from version control.

When implementing an MCP server, frameworks like `FastMCP` in the Python SDK abstract away much of the protocol’s complexity. Developers can define tools and resources using simple function decorators (`@mcp.tool()`, `@mcp.resource()`), with the framework automatically generating the necessary protocol-compliant definitions from Python type hints and docstrings. For local development and

testing, the `stdio` transport is highly efficient, allowing a host like Claude Desktop to launch the server as a child process and communicate directly. For production or multi-threaded scenarios, the HTTP+SSE transport is often more suitable.

Performance optimization is critical for production deployments. This involves employing asynchronous programming paradigms ( `asyncio` in Python) to handle concurrent requests efficiently, implementing intelligent multi-tier caching strategies to reduce latency and backend load, and designing for horizontal scalability with multiple server instances behind a load balancer. Context management itself is a key design pattern. Best practices include defining clear and concise context schemas, keeping context objects immutable where possible to prevent side effects, and implementing robust resource management to avoid memory leaks, especially in long-running sessions.

Above all, security must be integrated throughout the development lifecycle. This includes building robust user consent and authorization flows, implementing strict input validation on all data received from clients or LLMs, running servers with the least privilege necessary, and using containerization or other sandboxing techniques to isolate untrusted code. Comprehensive logging of all tool invocations and model decisions is essential for monitoring, debugging, and security auditing. By following these guidelines, developers can build MCP servers and clients that are not only functional and performant but also secure and trustworthy.

## 5. Comparative Analysis

The Model Context Protocol's design and purpose are best understood through comparison with both traditional communication protocols and other emerging standards for AI agent interoperability. Its introduction marks a significant evolution from established paradigms, tailored specifically to the dynamic, context-dependent nature of modern AI systems.

Compared to traditional APIs like **REST**, MCP offers a fundamental shift from rigid, predefined endpoints to dynamic, self-describing tools. In a RESTful architecture, adding a new capability or changing a parameter requires modifying the API contract, often leading to versioning complexities (e.g., `/v2/`) and breaking changes for clients. MCP, by contrast, allows servers to dynamically advertise their capabilities. Clients can discover and adapt to new tools or parameter changes at runtime without requiring code updates, fostering greater flexibility and future-proofing. Furthermore, while REST is inherently stateless, MCP is designed for stateful, session-oriented interactions, enabling it to maintain and propagate context across multiple turns of a conversation, which is crucial for complex agentic workflows.

While protocols like **GraphQL** address some of REST's data-fetching inefficiencies by allowing clients to request precisely the data they need, MCP goes further by standardizing not just data retrieval but also action execution (Tools) and instructional guidance (Prompts). High-performance protocols like **gRPC** offer efficient, type-safe communication ideal for microservices, but they do not inherently provide the AI-centric features of context management and dynamic tool discovery that are at the core of MCP.

MCP is also situated within a landscape of new agent interoperability protocols. Its development can be seen as an evolution from early academic efforts like **KQML** and **FIPA-ACL**, which were powerful but too complex for widespread adoption. It also builds upon the concepts of Retrieval-Augmented Generation (RAG) and function-calling APIs popularized by OpenAI. While function calling provided a structured way for LLMs to invoke external tools, it still required fragmented, manual integration for each service. MCP standardizes this interaction layer. Other contemporary protocols address different facets of agent communication: the **Agent Communication Protocol (ACP)** focuses on REST-native, multimodal messaging; the **Agent-to-Agent Protocol (A2A)** enables peer-to-peer task delegation;

and the **Agent Network Protocol (ANP)** facilitates open-network discovery using decentralized identifiers. A proposed adoption roadmap suggests using these protocols in concert, with MCP serving as the foundational layer for secure tool access and context ingestion within a broader, interoperable agent ecosystem.

## 6. Academic Perspectives and Future Directions

---

The academic community has begun to analyze the Model Context Protocol, recognizing it as a pivotal technology that addresses several long-standing challenges in the practical application of Large Language Models while also introducing new areas for research. Scholarly work positions MCP as a critical enabler for secure, adaptive, and scalable agentic workflows, moving beyond the limitations of stateless integration interfaces and ad-hoc security controls.

Research highlights MCP's effectiveness in mitigating critical LLM performance issues. One major challenge is **context overload**, where models are inundated with excessive or irrelevant information, degrading performance and increasing latency. MCP's architectural separation allows for strategic context filtering and prioritization, with its context management algorithms filtering noise and focusing the model on high-value data. Another challenge is the **token window limitation** of LLMs, which is often insufficient for processing extensive real-world datasets. MCP addresses this through intelligent context summarization, hierarchical context representation, and progressive loading techniques, where high-level summaries are provided first and details are loaded only as needed. The protocol also tackles the need for **dynamic context**, introducing mechanisms for continuous, differential updates that allow models to adapt to real-time changes in an environment without costly reprocessing.

Despite these advantages, academic analysis also points to open challenges and a rich set of future research directions. Security remains a primary concern, with researchers actively investigating mitigation strategies for threats such as cross-server privilege escalation, persistent-context tampering, installer spoofing, and tool name conflicts. The development of comprehensive security frameworks, including cryptographic hardening and capability-graph isolation, is a key area of focus.

Future research directions proposed in the literature are broad. They include protocol-level extensions to support multimodal payloads and more efficient transport mechanisms like QUIC. There is a call for simplifying MCP server deployment, potentially through no-code interfaces, to democratize AI automation. Further work is needed on performance optimization, enhancing model interpretability within MCP frameworks, and establishing robust governance models for the growing ecosystem. Architectural innovations, such as hybrid systems that combine rule-based automation with LLM-driven logic, are also being explored. This body of research underscores MCP's significance while mapping a clear path for its continued evolution and maturation.

## 7. Conclusion

---

The Model Context Protocol represents a foundational shift in the architecture of AI applications, moving the industry from a fragmented landscape of custom integrations to a standardized, interoperable ecosystem. By providing a universal language for AI models to securely access and interact with external tools and data, MCP directly addresses the critical barriers of scalability, security, and contextual awareness that have constrained the real-world utility of even the most advanced LLMs. Its well-defined host-client-server architecture, coupled with a flexible set of primitives like Tools, Resources, and Prompts, provides a robust framework for building the next generation of sophisticated, agentic systems.

However, the protocol's power is matched by its potential for misuse. The comprehensive security analysis reveals a complex threat landscape, from prompt-injection-based tool poisoning to classic server-side vulnerabilities. The future security and trustworthiness of the MCP ecosystem depend not only on the continued evolution of the protocol itself—particularly the standardization of authentication and authorization—but also on the diligent adoption of security-first best practices by the entire community of developers.

For software architects, AI engineers, and technical decision-makers, MCP is not merely another protocol to be learned; it is a strategic technology that will shape the future of AI development. Evaluating and implementing MCP requires a deep understanding of its architecture, a vigilant approach to its security implications, and a commitment to best practices. As the protocol matures and its ecosystem of tools and integrations expands, it is poised to become an indispensable component of the modern AI technology stack, enabling the creation of intelligent systems that are more connected, capable, and contextually grounded than ever before.

## 8. References

---

[A Survey of the Model Context Protocol \(MCP\): Standardizing Context to Enhance Large Language Models \(LLMs\)](https://www.preprints.org/manuscript/202504.0245/v1) (https://www.preprints.org/manuscript/202504.0245/v1)

[A survey of agent interoperability protocols: Model Context Protocol \(MCP\), Agent Communication Protocol \(ACP\), Agent-to-Agent Protocol \(A2A\), and Agent Network Protocol \(ANP\)](https://www.researchgate.net/publication/391461179_A_survey_of_agent_interoperability_protocols_Model_Context_Protocol_MCP_Agent_Communication_Protocol_ACP_Agent-to-Agent_Protocol_A2A_and_Agent_Network_Protocol_ANP) (https://www.researchgate.net/publication/391461179\_A\_survey\_of\_agent\_interoperability\_protocols\_Model\_Context\_Protocol\_MCP\_Agent\_Communication\_Protocol\_ACP\_Agent-to-Agent\_Protocol\_A2A\_and\_Agent\_Network\_Protocol\_ANP)

[A survey of agent interoperability protocols: Model Context Protocol \(MCP\), Agent Communication Protocol \(ACP\), Agent-to-Agent Protocol \(A2A\), and Agent Network Protocol \(ANP\)](https://arxiv.org/abs/2505.02279) (https://arxiv.org/abs/2505.02279)

[A Survey on Model Context Protocol: Architecture, State-of-the-art, Challenges and Future Directions](https://www.techrxiv.org/users/913189/articles/1286748-a-survey-on-model-context-protocol-architecture-state-of-the-art-challenges-and-future-directions) (https://www.techrxiv.org/users/913189/articles/1286748-a-survey-on-model-context-protocol-architecture-state-of-the-art-challenges-and-future-directions)

[AI Model Context Protocol \(MCP\) and Security](https://community.cisco.com/t5/security-blogs/ai-model-context-protocol-mcp-and-security/ba-p/5274394) (https://community.cisco.com/t5/security-blogs/ai-model-context-protocol-mcp-and-security/ba-p/5274394)

[A Primer on Model Context Protocol \(MCP\) Secure Implementation](https://cloudsecurityalliance.org/blog/2025/06/23/a-primer-on-model-context-protocol-mcp-secure-implementation) (https://cloudsecurityalliance.org/blog/2025/06/23/a-primer-on-model-context-protocol-mcp-secure-implementation)

[How to Implement Model Context Protocol](https://astconsulting.in/model-context-protocol/how-to-implement-model-context-protocol) (https://astconsulting.in/model-context-protocol/how-to-implement-model-context-protocol)

[Introducing the Model Context Protocol](https://www.anthropic.com/news/model-context-protocol) (https://www.anthropic.com/news/model-context-protocol)

[Is Your AI Safe? Threat Analysis of MCP \(Model Context Protocol\)](https://www.cyberark.com/resources/threat-research-blog/is-your-ai-safe-threat-analysis-of-mcp-model-context-protocol) (https://www.cyberark.com/resources/threat-research-blog/is-your-ai-safe-threat-analysis-of-mcp-model-context-protocol)

[Master Model Context Protocol \(MCP\): A Tutorial & Guide](https://content.trickle.so/blog/master-model-context-protocol-tutorial-guide) (https://content.trickle.so/blog/master-model-context-protocol-tutorial-guide)

[MCP Security Exposed: What You Need to Know Now](https://live.paloaltonetworks.com/t5/community-blogs/mcp-security-exposed-what-you-need-to-know-now/ba-p/1227143) (https://live.paloaltonetworks.com/t5/community-blogs/mcp-security-exposed-what-you-need-to-know-now/ba-p/1227143)

[MCP Security: How to Secure Your AI Agent Toolchain](https://stytech.com/blog/mcp-security/) (https://stytech.com/blog/mcp-security/)

[MCP Security: Insights from Thinkata](https://thinkata.com/news/insights/mcp-security) (https://thinkata.com/news/insights/mcp-security)

[Model Context Protocol \(MCP\) - Landscape, Security Threats and Future Research Directions](https://www.scribd.com/document/848536652/Model-Context-Protocol-MCP-Landscape-Security-Threatsand-Future-Research-Directions) (https://www.scribd.com/document/848536652/Model-Context-Protocol-MCP-Landscape-Security-Threatsand-Future-Research-Directions)

[Model Context Protocol \(MCP\): Foundation for AI or a Looming Risk? — AI Innovations and Insights 37](https://pub.towardsai.net/model-context-protocol-mcp-foundation-for-ai-or-a-looming-risk-ai-innovations-and-insights-37-2b8dad940c69) (https://pub.towardsai.net/model-context-protocol-mcp-foundation-for-ai-or-a-looming-risk-ai-innovations-and-insights-37-2b8dad940c69)

[Model Context Protocol \(MCP\): Landscape, Security Threats, and Future Research Directions](https://arxiv.org/abs/2503.23278) (https://arxiv.org/abs/2503.23278)

[Model Context Protocol \(MCP\): The Complete Tutorial](https://www.claudemcp.com/docs) (https://www.claudemcp.com/docs)

[Model Context Protocol - A Comprehensive Guide](https://meghashyamthiruvedula.medium.com/model-context-protocol-a-comprehensive-guide-41c8b56a61f3) (https://meghashyamthiruvedula.medium.com/model-context-protocol-a-comprehensive-guide-41c8b56a61f3)

[Model Context Protocol Documentation](https://www.firemcp.com/documentation) (https://www.firemcp.com/documentation)

[Model Context Protocol: Enhancing LLM Performance for Observability and Analytics](https://ejournals.org/ejcsit/wp-content/uploads/sites/21/2025/05/Model-Context-Protocol.pdf) (https://ejournals.org/ejcsit/wp-content/uploads/sites/21/2025/05/Model-Context-Protocol.pdf)

[Model Context Protocol Introduction](https://modelcontextprotocol.io/introduction) (https://modelcontextprotocol.io/introduction)

[Model Context Protocol Introduction](https://modelcontextprotocol.io/introduction) (https://modelcontextprotocol.io/introduction)

[Model Context Protocol on GitHub](https://github.com/modelcontextprotocol) (https://github.com/modelcontextprotocol)

[Model Context Protocol Server Setup Guide: Step-by-Step Tutorial](https://www.byteplus.com/en/topic/541333?title=model-context-protocol-server-setup-guide-step-by-step-tutorial) (https://www.byteplus.com/en/topic/541333?title=model-context-protocol-server-setup-guide-step-by-step-tutorial)

[Model Context Protocol Servers: A Novel Paradigm for AI-Driven Workflow Automation and Comparative Analysis with Legacy Systems](https://www.researchgate.net/publication/389687667_Model_Context_Protocol_Servers_A_Novel_Paradigm_for_AI-Driven_Workflow_Automation_and_Comparative_Analysis_with_Legacy_Systems) (https://www.researchgate.net/publication/389687667\_Model\_Context\_Protocol\_Servers\_A\_Novel\_Paradigm\_for\_AI-Driven\_Workflow\_Automation\_and\_Comparative\_Analysis\_with\_Legacy\_Systems)

[Model Context Protocol Specification](https://modelcontextprotocol.info/specification/) (https://modelcontextprotocol.info/specification/)

[Model Context Protocol Specification 2025-06-18](https://modelcontextprotocol.io/specification/2025-06-18) (https://modelcontextprotocol.io/specification/2025-06-18)

[Model Context Protocol Specification: Security Best Practices](https://modelcontextprotocol.io/specification/draft/basic/security_best_practices) (https://modelcontextprotocol.io/specification/draft/basic/security\_best\_practices)

[Model Context Protocol Tutorials](https://modelcontextprotocol.io/tutorials/building-mcp-with-llms) (https://modelcontextprotocol.io/tutorials/building-mcp-with-llms)

[Model Context Protocol Tutorials](https://modelcontextprotocol.info/docs/tutorials/) (https://modelcontextprotocol.info/docs/tutorials/)

[Model Context Protocol Tutorial \(DataCamp\)](https://www.datacamp.com/tutorial/mcp-model-context-protocol) (https://www.datacamp.com/tutorial/mcp-model-context-protocol)

[The Model Context Protocol \(MCP\): A Complete Tutorial](https://medium.com/@nimritakoul01/the-model-context-protocol-mcp-a-complete-tutorial-a3abe8a7f4ef) (https://medium.com/@nimritakoul01/the-model-context-protocol-mcp-a-complete-tutorial-a3abe8a7f4ef)

[The Security Risks of Model Context Protocol \(MCP\)](https://www.pillar.security/blog/the-security-risks-of-model-context-protocol-mcp) (<https://www.pillar.security/blog/the-security-risks-of-model-context-protocol-mcp>)

[Understanding and mitigating security risks in MCP implementations](https://techcommunity.microsoft.com/blog/microsoft-security-blog/understanding-and-mitigating-security-risks-in-mcp-implementations/4404667) (<https://techcommunity.microsoft.com/blog/microsoft-security-blog/understanding-and-mitigating-security-risks-in-mcp-implementations/4404667>)

[Understanding Model Context Protocol: A Practical Guide for Intermediate AI Developers](https://mskadu.medium.com/understanding-model-context-protocol-a-practical-guide-for-intermediate-ai-developers-877375757d50) (<https://mskadu.medium.com/understanding-model-context-protocol-a-practical-guide-for-intermediate-ai-developers-877375757d50>)

[Uncovering MCP Security](https://www.querypie.com/resources/discover/white-paper/18/uncovering-mcp-security) (<https://www.querypie.com/resources/discover/white-paper/18/uncovering-mcp-security>)

[Untitled](https://arxiv.org/abs/2504.08623) (<https://arxiv.org/abs/2504.08623>)

[Untitled](https://docs.anthropic.com/en/docs/agents-and-tools/mcp) (<https://docs.anthropic.com/en/docs/agents-and-tools/mcp>)